

TO: SwiftCart Analytics Review Board
FROM: Swamini Wagh (2025TT12078)
DATE: June 2025
SUBJECT: The Cost of Being Wrong — Asymmetric Inventory Optimization Results

1. Total Business Cost Comparison

Scoring formula applied across the unseen test period (Oct–Dec 2023):

Total Business Cost = (Missed Stockouts × \$10) + (Units Spoiled × \$1)

Model	Missed Stockouts (FN)	Total Business Cost
Baseline — 50% threshold	288	\$12,473
Cost-Sensitive (class_weight 1:10)	290	\$12,493
Tuned Threshold — 9.09% (Best)	38	\$9,973
XGBoost Custom Loss + Tuned Threshold	100	\$10,593

The tuned threshold model delivered a **\$2,500 cost reduction (20.0%)** over the baseline by reducing missed stockout events from 288 to 38 on unseen data.

2. Mathematical Approach – Three asymmetric optimization techniques were implemented:

A) Cost-Sensitive Class Weights (Random Forest)

`class_weight = {0: 1, 1: 10}` — the algorithm multiplies the internal loss on every missed stockout by 10x, forcing it to prioritize stockout recall over overall accuracy.

B) Decision Threshold Tuning

The optimal threshold was derived mathematically from the cost ratio:

$P^* = C_{FN} / (C_{FN} + C_{FP}) = 10 / (10 + 1) = 9.09\%$

Instead of the default 50% confidence requirement, a restock alert fires at just 9.09% predicted probability — aggressively protecting against the more expensive error class.

C) Custom Asymmetric Loss Function (XGBoost)

A custom gradient/hessian objective was written where:

`weight = 10.0 if stockout, else 1.0`

`grad = weight × (sigmoid( $\hat{y}$ ) -  $y$ )`

Every boosting round penalizes stockout misses 10× harder than spoilage errors, steering the model toward commercial efficiency rather than raw accuracy.

3. Spoilage Constraint Check

Metric	Value
Total units ordered (Oct–Dec test period)	137,670
Total units spoiled (Oct–Dec test period)	9,593
Network-wide spoilage rate	6.97%
Constraint ceiling	5.00%

Baseline statement: The observed spoilage rate of 6.97% exceeds the 5% ceiling. However, this spoilage is a fixed property of the inventory dataset - it reflects procurement decisions already made before the model runs. The classification model predicts *whether* a stockout will occur; it does not directly control reorder quantities or spoilage volumes. In a production deployment, the asymmetric model would be paired with a reorder quantity optimizer to bring spoilage under the 5% ceiling. The model's primary contribution: reducing missed stockouts by 87% (288 → 38) - directly addresses the higher-cost failure mode of the two.